

SWIPE →

5 guardrails your AI agents need before they touch anything real.


STEP 01 / 05

Scope it

Write down exactly which systems, accounts, and data the agent can reach. Everything else is off-limits by default.


04


STEP 02 / 05

Give it a stop rule

“If you’re blocked, stop and ask a human.” A locked door is a signal — not a puzzle to solve.

02




STEP 03 / 05

Log everything

Every action, file, and login attempt —
stored somewhere a human actually
reviews.


03


STEP 04 / 05

Pre-write your disclosure plan

Who gets told, how, and within how many hours. Not a public inbox three months later.

04



STEP 05 / 05

Name an owner

One accountable human for every agent running in production.

05

FREE • 1-PAGE RESOURCE

Get the Agent Guardrails Checklist.

- Access questions to answer before go-live
- Stop rules worth hard-coding
- What to log — and who reviews it
- A fill-in-the-blank disclosure plan

Comment

AGENT

and we'll send it. [Follow for weekly AI briefs.](#)