



It hit a locked door. So it found another way.

What an OpenAI research agent did on a government health portal, and 5 guardrails so yours doesn't.

Jason Redfoot • Founder, Pragentic.ai

- 01 Scope it
- 02 Give it a stop rule
- 03 Log everything
- 04 Pre-write your disclosure plan
- 05 Name an owner

SWIPE →

WHAT HAPPENED

~3 months

In June, an OpenAI research agent accessed non-public files on an Australian government health statistics portal. The government was notified Sept 10, via a public mailbox.

No personal data appears to have been accessed. Three more government sites are now being checked.

It wasn't malicious. It was persistent.

Now put that agent in a 20-person firm.
Access to your client portal, shared drive, and
CRM. Told to “pull the numbers.” Blocked.

What does it do next? Most firms never wrote
it down.

01

Scope it

List the exact systems, accounts, and data the agent can reach. Everything else is off-limits by default.

Try this: one list per agent, kept where the owner can see it.



02

Give it a stop rule

“If you’re blocked, stop and ask a human.” A locked door is a signal, not a puzzle to solve.

Try this: put the rule in the agent’s instructions, word for word.




03

Log everything

Every action, file, and login attempt, stored somewhere a human actually reviews.

Try this: a weekly 15-minute log review on the owner's calendar.




04

Pre-write your disclosure plan

Who gets told, how, and within how many hours.

Not a public mailbox three months later.



05

Name an owner

One accountable human for every agent running in production.

If nobody owns it, nobody stops it.

Before any agent goes live:

- ☐ Scoped
- ☐ Stop rule
- ☐ Logged
- ☐ Disclosure plan
- ☐ Owner named

FREE · CHECKLIST

Want the full Agent Guardrails Checklist? Comment

AGENT

on this post and I'll send it. [Follow Jason Redfoot](#) for weekly AI briefs.

Save this for your next AI rollout.